

# Credit Default Risk

Predicting how likely each applicant  
is of repaying a loan?



Goal, identify if a new client shows a high risk for loan default.

## How can this help?



Reduce Uncertainty



Proportional  
Disbursement



Risk Reduction

**Doesn't leave business on the table!**

# Cleaning the data first involves understanding the columns (add infos)

There are 63 clumns with different data types:

- 30 floating point numbers.
- 22 integer numbers – some are numerical categories.
- 11 Strings – Categories involved.

---

Referral\_code has the highest category counts with 7805 categories.

```
preferred_contact      3
referral_code          7805
account_status_code    5
employment_type        16
education              5
marital_status         3
state                  20
loan_type              12
loan_purpose            8
origination_channel    4
marketing_campaign      26
dtype: int64
```




# Cleaning steps

Since the ML model can't inherently deal with text, the data must be converted to appropriate numbers. Any significant distortion/noise in the model must be removed as much as possible.

- Converting string categorical columns into numerical – Label encoding.
- Converting string categorical columns into numerical and adding new columns to indicate the presence of categorical variables – One hot encoding.
- Replacing illogical outliers with empty values (NAN values).
- Imputing empty cells with the median of the values. In some cases, imputation is approached with a certain grouping.
- Dealing with a few anomalies.
- Changing invalid entries into valid entries.

# Predictive Modeling – Outcome of the model is expected to identify the potential that someone will default on a loan.




Expected Target Outcome: 0 or 1, 0 – Not a defaulter, 1 – potential defaulter.

Performance Metrics used : Accuracy.

Models currently used : Logistic regression, Random forest, XGBoost.

# Training and Testing datasets were subjected to the same feature engineering to evaluate the model.



- Out of the main training dataset, a certain percentage is kept untrained to test the model's performance.
- Training set and validation set are split in following percentages: 80% : 20%.
- On the testing set, the target labels are hidden, until the performance is evaluated.

# Logistic regression

The accuracy: 0.7387222222222222

Roc AUC score: 0.7238519289675215

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	0.98	0.74	0.84	17081
---	------	------	------	-------

1	0.13	0.71	0.22	919
---	------	------	------	-----


accuracy			0.74	18000
----------	--	--	------	-------

macro avg	0.55	0.72	0.53	18000
-----------	------	------	------	-------

weighted avg	0.94	0.74	0.81	18000
--------------	------	------	------	-------

Confusion matrix [[12647 4434]  
[ 269 650]]

# Logistic regression



# Random Forest


The accuracy: 0.9437222222222222

Roc AUC score: 0.5492429370166687

	precision	recall	f1-score	support
0	0.95	0.99	0.97	17081
1	0.34	0.11	0.17	919
accuracy			0.94	18000
macro avg	0.65	0.55	0.57	18000
weighted avg	0.92	0.94	0.93	18000

Confusion matrix [[16886 195]  
[ 818 101]]

# Logistic regression



# XGBoost

The accuracy: 0.8981666666666667


Roc AUC score: 0.6652645377376526

	precision	recall	f1-score	support
0	0.97	0.92	0.95	17081
1	0.22	0.41	0.29	919
accuracy			0.90	18000
macro avg	0.60	0.67	0.62	18000
weighted avg	0.93	0.90	0.91	18000

Confusion matrix [[15794 1287]

[ 546 373]]

# XGBoost



# Currently there are 3 models used.

Model Name	Hyperparameters	Accuracy
LogisticRegression	(C=2, class_weight=None, dual=False, fit_intercept=True, intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=1, penalty='l2', random_state=None, solver='liblinear', tol=0.0001, verbose=0, warm_start=False)	74%
RandomForestClassifier	(bootstrap=True, class_weight=None, criterion='gini', max_depth=None, max_features='auto', max_leaf_nodes=None, min_impurity_decrease=0.0, min_impurity_split=None, min_samples_leaf=1, min_samples_split=2, min_weight_fraction_leaf=0.0, n_estimators=500, n_jobs=-1, oob_score=False, random_state=50, verbose=1, warm_start=False)	94%
XGBClassifier	(base_score=0.5, colsample_bylevel=1, colsample_bytree=1, gamma=0, learning_rate=0.1, max_delta_step=0, max_depth=5, min_child_weight=1, missing=None, n_estimators=250, nthread=-1, objective='binary:logistic', reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=0, silent=True, subsample=1)	90%



# Random Forest is currently the best chosen model, Until further work.

- More feature engineering could be done, such as mathematical transformation, possibly: certain columns to log and such.
- Advanced techniques like SMOTE could be deployed to handle the class imbalance problems.
- DNN's can be used.
- Long short term memory networks could be used to incorporate time series data.
- Advanced forms of stacking can be used apart from voting.
- Recursive feature selection can be used to reduce the features.



Questions?

Thank you!